

ANIRBAN GOSWAMI

Data Engineer | Databricks | PySpark | Azure

+91 8927580276 • anirbang324@gmail.com • [LinkedIn](#) • [GitHub](#) • Kolkata, West Bengal

PROFESSIONAL SUMMARY

Results-driven Data Engineer with 2+ years at Accenture specializing in Databricks, PySpark, and cloud-based ETL pipelines across Azure, AWS, and GCP. Proven track record delivering high-impact migrations (700+ views, 20 TB+ data) and building production-grade data platforms.

EXPERIENCE

Data Engineering, Management & Governance Analyst | Accenture

Sept 2025 – Present | Kolkata

- Architected a PySpark accelerator on Databricks to automate migration of 700+ Hive views to Unity Catalog, reducing manual effort by 90% and saving 20+ hrs/week.
- Led centralization of metadata for 10+ datasets across business units, enforcing Unity Catalog governance standards; managed migration of 5+ enterprise apps spanning 1,000+ database objects.
- Collaborated with architects to define Unity Catalog permission models (row filters, column masks) ensuring least-privilege access across migrated tables.

Data Engineering, Management & Governance Associate | Accenture

Feb 2024 – Sept 2025 | Kolkata

- Built and optimized 15+ production Databricks ETL pipelines processing TBs of data daily, achieving a 40% reduction in runtime through partition tuning and AQE.
- Developed a cross-cloud migration tool (AWS S3 → GCP Cloud Storage) transferring 20 TB+ with 99.9% reliability using PySpark and checkpointing.
- Implemented Delta Lake best practices (OPTIMIZE, ZORDER, CDF) and automated data quality checks with Great Expectations, cutting downstream data defects by 60%.

Apprentice – Full-Stack Developer | LTIMindtree

May 2023 – Sept 2023 | Remote

- Contributed to 3+ full-stack projects using TypeScript, AngularJS, Core Java, and SQL, serving 500+ internal users.

PROJECTS

Real-Time E-Commerce Analytics Platform

Tech Stack: Databricks, PySpark Structured Streaming, Delta Lake, Kafka, ADF, Power BI, Unity Catalog

- Designed a Medallion Architecture (Bronze→Silver→Gold) ingesting real-time Kafka clickstream + nightly batch files, achieving sub-5-second latency for 50K+ events/min via Auto Loader and Structured Streaming.
- Implemented CDC (SCD Type 2) via APPLY CHANGES INTO, applied Liquid Clustering & Deletion Vectors (55% faster queries, GDPR purge support), and secured PII with Unity Catalog row filters/column masks.

Unity Catalog View Migration Accelerator

Tech Stack: Databricks, PySpark, Delta Lake, Unity Catalog, SQL

- Automated migration of 700+ Hive metastore views to Unity Catalog via a PySpark framework with dependency resolution and rollback, cutting migration errors by 70% and saving 20+ hrs/week.

SKILLS

Languages: Python, SQL, Java

Data Engineering: PySpark, Databricks (Lakeflow, Unity Catalog, Auto Loader), Delta Lake, Kafka, Spark Structured Streaming

Cloud & Storage: Azure (ADF, ADLS Gen2, Synapse), GCP (BigQuery, Cloud Storage, Dataflow), AWS S3, CloudWatch

Platforms & DBs: Microsoft Fabric, Power BI, Great Expectations, MySQL, PostgreSQL, MongoDB

Web / Misc & DevOps: JavaScript, TypeScript, Angular, Django, NumPy, Pandas, Scikit-learn, Git, GitHub, GitLab, CI/CD

CERTIFICATIONS

Databricks Certified Data Engineer Professional (2026) • Databricks Certified Data Engineer Associate (2025) • Claude Certified Architect – Foundations, Anthropic (2026) • Azure Fundamentals AZ-900 (2024) • Azure Data Fundamentals DP-900 (2024) • GCP Associate Cloud Engineer (2025)

EDUCATION

B.Tech in Computer Science, Asansol Engineering College — CGPA 8.9/10

2019 – 2022

Diploma in Computer Science & Technology, Dr. B.C. Roy Polytechnic — CGPA 8.0/10

2016 – 2019

ACHIEVEMENTS & EXTRACURRICULAR

- ACE Award (Accenture) – Delivered a critical client data migration project 2 weeks ahead of schedule, improving customer satisfaction score by 15%.
- Solved 200+ Data Structures & Algorithms problems on LeetCode and Codeforces.